

Oct 30, 2013
Jeremy Li

Evaluation of HP3PAR 7400

HP introduced a newer model of [3PAR StoreServ 7000 Series](#) for Enterprise Storage: 7200, 7400 and 7450 arrays in November, 2012.

According to the [DCIG 2014 ENTERPRISE MIDRANGE ARRAY BUYER'S GUIDE](#), HP earned the "Best-In-Class" Ranking among enterprise midrange arrays, while EMC VNX earned the "Recommended" Ranking.

HP is a [Fortune 500 company](#) and has nearly 350,000 employees globally and delivered the HP 3PAR StoreServ 7000 Series in December, 2012.

HP 3PAR 7400 is the best new technology in the storage market today because it was designed for today's virtualized environments and is optimized for cloud deployments. However, ease of use, performance, and cost savings features are valuable to all customers. Below are a few highlights:


- 3PAR Thin Provisioning - Industry leading technology to maximize storage utilization
- 3PAR Autonomic Storage Tiering – Automatically optimizes using multiple classes of storage
- 3PAR Virtual Domains – Multi-tenancy for service providers and private clouds
- 3PAR Dynamic Optimization – Workload management and load balancing
- 3PAR Mesh-Architecture – Advanced shared-memory architecture
- A single array can be scaled out to 4 mesh active/active hardware-based 4th generation ASIC (Figure 2) [controllers](#) up to 864TB, and soon 1.2PB
- Thin Built in Zero Detect – It is almost like inline dedupe. That's why HP guarantees its customers a 50% storage savings
- The starting list price is under \$45,000 due to its famous Massively Parallel Architecture (MPA) – using small form factor (SFF) SAS drives 10K RPM that can outperform its competitors' large form factor (LFF) SAS drives 15K RPM that is more expensive than SFF SAS drives at 10K RPM

- Virtual Volumes, as illustrated in the Figure 1, along with Autonomical Wide-striping across Logical Disks (LDs), also known as “Disk Striping” at 128KB block size across all available disks (HDDs), are the keys to form a massively parallel architecture from HP 3PAR, which automatically spreads all workloads over all internal disks resources. As a result, the MPA can deliver higher and more predictable levels of performance with high IOPS, low latencies and full capacity utilization.
- The MPA can give customers higher VMs per Host ratio. For example, HP guarantees customers: 20:1 ratio on its HP 3PAR 7000, targeted for a mid-range storage customers with tier 1 capability usually from EMC vMAX storage.
- Integrated XOR Engine that creates parity information quickly under the hardware-based acceleration - [XOR](#) is used in RAID 3–6 for creating parity information. For example, RAID can "back up" bytes 10011100 and 01101100 from two (or more) hard drives by XORing the just mentioned bytes, resulting in (11110000) and writing it to another drive. Under this method, if any one of the three hard drives is lost, the lost byte can be re-created by XORing bytes from the remaining drives. For instance, if the drive containing 01101100 is lost, 10011100 and 11110000 can be XORed to recover the lost byte. (Source: Wikipedia).
- Tightly VMware integration because the VMware ESXi OS was developed on the 3PAR storage from the very beginning.

Architectural Attributes

Benefits for VMware include:

1. VMware application service levels are higher and more predictable
2. More Virtual Servers can be consolidated per ESX Server, enhancing VMware ROI
3. More transaction-intensive applications can be supported in a virtualized ESX environment
4. Fewer arrays and less capacity are required to support a given VMware deployment



Example: 4-Node 7400 with 8 drive enclosures

Process step	Phase state
<ul style="list-style-type: none">Physical drives (PD) are automatically formatted in 1GB chunklets	Disk initialization
<ul style="list-style-type: none">Chunklets are bound together to form Logical Disks (LD) in the format defined in CPGs	Defines RAID level, step size, set size and redundancy
<ul style="list-style-type: none">Virtual Volume are built on LDs across all nodes	Autonomical wide-striping across LDs
<ul style="list-style-type: none">Virtual Volumes can now be exported as LUNs to servers	Present LUNs ready to go

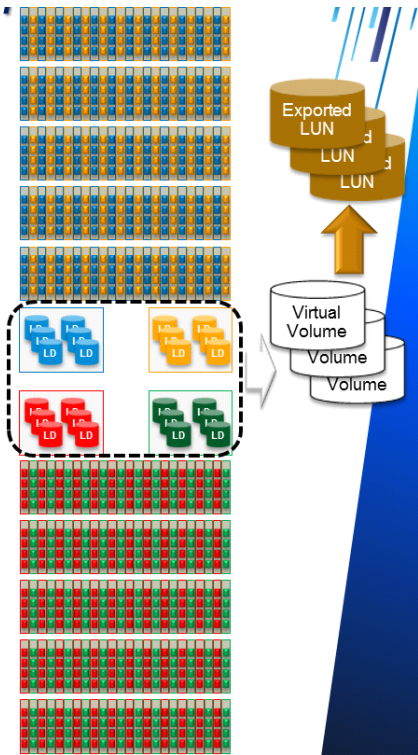


Figure 1 Autonomical wide-striping across all available logical disks

HP 3PAR Gen4 ASIC

Hardware Based for Performance

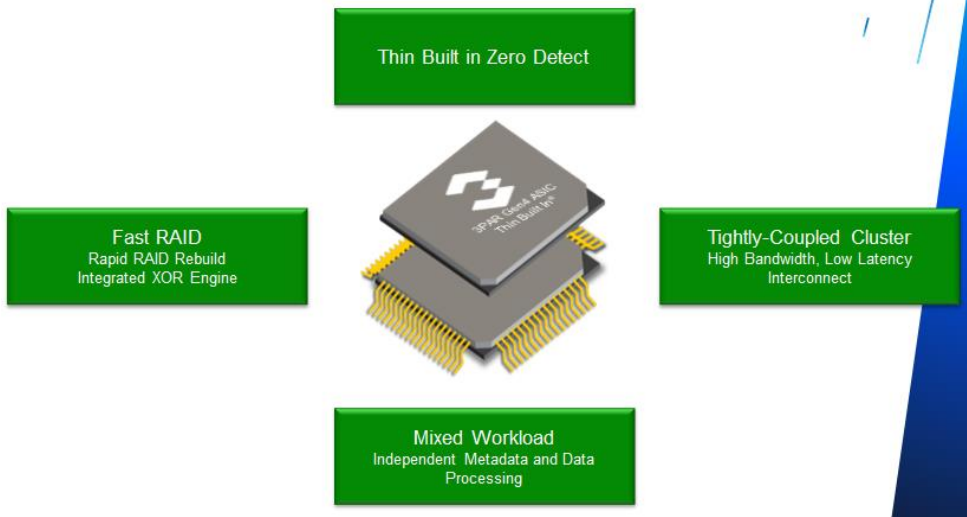


Figure 2 HP 3PAR 4th Generation ASIC

That's why 6 to 8 out of top 10 hosting companies choose HP 3PAR.

Disk Striping and Virtual Volume

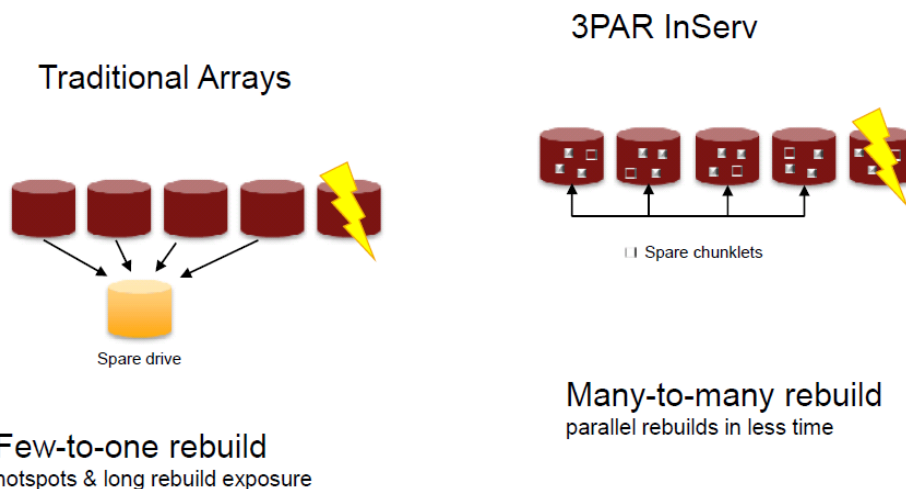
Most storage vendors use Disk Striping technique in order to boost the storage performance. However, this technique is not good enough in today's storage.

The 3PAR MPA delivers the great storage performance and reduces a failed drive recovery time from days to hours. As of today, HP 3PAR, Dell Compellent and IBM/NetApp E-Series storage with DDP enable, are only a few vendors who can recover a failed disk from days to hours with a 3TB HDD at the 80% capacity full. The more HDDs in the HP 3PAR system, the lesser recovery time will be because 3PAR system spreads all workloads over all internal disks due to its MPA regardless whether a RAID level with 16 drives or 24 drives is configured. This is because the exported LUNs are based on a Virtual Volume that is built on Logic Disks across all nodes detailed in Figure 1 above.

The Rebuilding Time from a Failed Drive

A typical rebuilding time from a failed drive will take from 3 days to 5 days in most storage vendors' system. IBM DCS3700 and DS3500 can reduce the rebuilding time from days to hours. [Click here](#) for details.

However, many other vendors' rebuilding time for a failed disk drive will be much longer since they do not have the capability of either DDP or spare chunklets from HP 3PAR or Dell Compellent. Below, a screenshot illustrates a traditional RAID rebuilding method vs. a parallel rebuilds method:



With the DDP or Chunklets method, a dedicated parity disk drive is no longer used to recover data from a failed drive. Therefore, a bottleneck due to a hot spot from a dedicated parity drive after a new drive is inserted into a storage system is eliminated.

Below are summary of HP 3PAR advantages over other Storage Vendors:

HP 3PAR from 7200 series to 10000 series can archive “Polymorphic Simplicity with One Architecture: One Operating System, One Interface and One Feature Set.

Click [here](#) and [there](#) for a comparison among HP3PAR, EMC VMAX and VNX (Source: HP). Good news: EMC is going to release ALL FLASH products based on XtremIO (XPECT MORE) to help enterprise(s) to increase the storage performance and IT resource utilization, in turn, reduce the TCO.

Hadoop

Hadoop is pioneered by Google’s [MapReduce](#), a programming model and framework where an application is broken down into numerous small parts.

[Apache Hadoop](#) is an [open-source software framework](#) for storing and large scale processing of data-sets on clusters of [commodity hardware](#) (Source: Wikipedia). Intel and other leading IT vendors are making Hadoop the next platform, as illustrated in Figure 3.

EMC Isilon might be a better choice for Hadoop due to its distributed computing architecture – [parallel processing](#) with rapid data transfer rates among nodes and still having the uptime in case of a node failure (N+1, N+2, N+3 and N+4) by OneFS scale-out operating system (OS) and its unique technique by using a pointer to generate three identical copies (a compliance for a Hadoop) on a scale-out storage in Big Data Analytics scenarios – **a massive storage savings only Isilon can deliver at this time** because

- At least three (3) identical data must be copied to the Hadoop cluster before analysis can begin
- Simple install, manage and scale without going through a forklift upgrade
- Delivers powerful yet simple solutions for enterprises that want to manage their data, not their storage.
- Overcome the [traditional “Share-Nothing” Hadoop](#) with [innovative “Share-Everything” Hadoop](#)

Below are highlights of EMC Isilon:

- Single File System with [Simplicity](#)
- [High Performance Computing](#)
- [Automated Tiering](#)
- [Easy Growth](#)
- [Linear Scalability](#)

Additional Information from [B. Scott Cassell](#), EMC²| Isilon Storage Division

My notes is my personal opinion and does not represent the views of the County of Los Angeles

- ONLY Isilon lets Hadoop users separate the storage component from the computer component. Meaning with Isilon, you can add just compute to MapReduce jobs faster without adding any additional storage, or just add storage to hold larger data sets without adding any additional compute components.
- ONLY Isilon lets Hadoop users start MapReducing immediately without the need to duplicate and move the duplicated data to the target data nodes.
- ONLY Isilon lets Hadoop users perform analytics-in-place because there is no need to duplicate and ship data. Analytics can be performed on the one copy of actual original data in real time without preparatory steps.

Note: Managing EMC VMAX storage requires a few highly skilled storage administrators

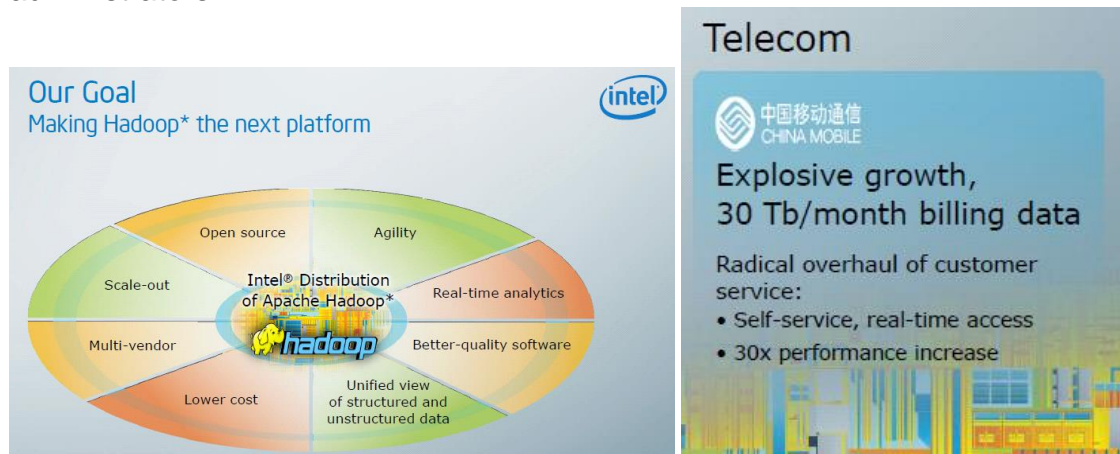


Figure 3 Intel is making Hadoop the next platform

Note: INTEL and CHINA MOBILE, the No. 1 brand name in China from 2010 to 2013, are working together and have achieved the savings of \$30 million per month via Hadoop technology, making Hadoop the next platform.

Caution:

EMC Isilon is not for everyone. EMC even puts Isilon in the midrange storage while VNX and VMAX should be considered at a high-end storage environment.

Below is one of the keynote slides at the EMC Forum, August, 2013 in Long Beach.

Best Of Breed Storage & Backup

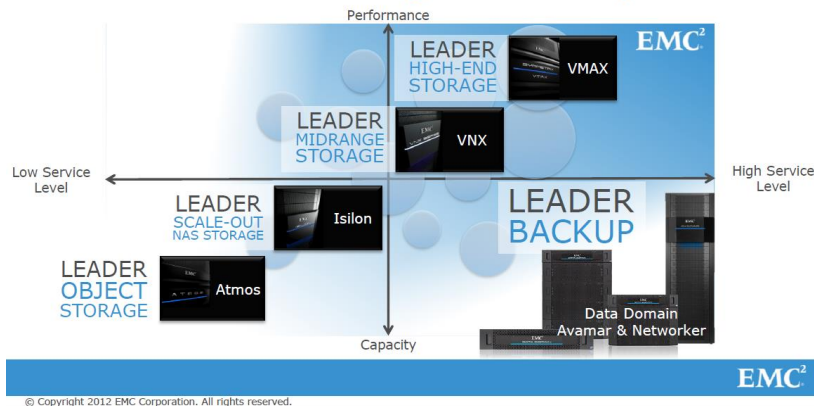


Figure 3 Isilon is for Midrange Storage

The reason is simple:

Neither Hadoop technology nor NAS only storage (no application can be run from NAS) must be used by an enterprise while a unified storage such as EMC VNX can do both in one box.

For a very large enterprise such as Facebook, EMC Isilon is an excellent choice because every box is inserted into its cluster, higher IOPS and more storage will be achieved with additional parallel processing capability (scale-out NAS) at the cost of additional CPU, Memory, Power Supplies, Infinite Band (IB) and 10/1 Gigabit connections. When a cluster becomes large, a massive parallel processing capability is formed.

Conclusion

HP 3PAR 7400 is the best new technology in the storage market today because it was designed for today's virtualized environments and is optimized for cloud deployments. However, ease of use, performance, and cost savings features are valuable to all customers. Below are a few highlight points:


- 3PAR Thin Provisioning - Industry leading technology to maximize storage utilization
- 3PAR Autonomic Storage Tiering – Automatically optimizes using multiple classes of storage
- 3PAR Virtual Domains – Multi-tenancy for service providers and private clouds
- 3PAR Dynamic Optimization – Workload management and load balancing
- 3PAR Mesh-Architecture – Advanced shared-memory architecture

The comparison of primary value to any enterprise(s) should be based on cost per terabyte (TB), IOPS, Density, Total Cost of Ownership (TCO) and Return on

Investment (ROI). The comparison will provide real financial benefit to the enterprise(s) over a 3-5-year period.

IT is about providing a specific level of service at the lowest possible cost. The value to the business is the ability to improve the service while saving money.

In order to increase application performance at lower cost, today's CIOs in any organization(s) are looking to consolidate their mission-critical applications on fewer, faster servers and to virtualize those applications to make more efficient utilization of their data center resources and often are being asked to do more with fewer resources, less time and less funding, and transform and guide an enterprise into a lean organization. Therefore, carefully examining its internal IT operations and efficiency, including using available technology in new efficient ways, often leads to trim many hidden costs and yield greater flexibility to focus on core business and pursue the innovation.



Gartner introduced the idea of "Total Cost of Ownership" (TCO) in 1986, as illustrated in the screenshot from left. The initial acquisition cost is only a part of the equation of TCO. More software purchased, more hidden cost arisen. Therefore, any company should select a solution in accordance with its own

environment needs.

Source: <http://www.alligatorsql.com/solutions/tco/index.jsp>

Below is an excerpt from [Navigator, THE CLIPPER GROUP](#)

"By reducing the number of physical servers populating the data center, the CIO can reduce the number of systems administrators required to drive the IT infrastructure, as well as reducing the amount of energy necessary to power the data center, and the amount of floor space required to house it. These last two points are especially critical as enterprise data centers approach maximum capacity in both of these categories. In fact, if either is exceeded, the enterprise may be forced to build out a brand new data center at a cost of millions of dollars."

Challenge


EMC is still the No. 1 disk storage market leader with highest market share globally due to its reputation and superior technical support. Choosing EMC will not let you go wrong, as long as the capital investment is not a question. Many enterprises would like to choose the incumbent leaders in addition to traditional stereotype.

However, [EMC stock](#) is at about \$24.00 for a long time and has been going down by 6.44% on a Year-to-Date basis due to the storage vendors' thin margin competition, while [HP stock](#) is at a similar value of EMC's for a long time.

Additional Information:

Any storage system will have a different usable capacity, even though the raw storage capacity is identical.

Below is the snapshot of a few vendors' usable capacity, snapshot usage, SATA Checksum overhead, File System Reserve and Parity.



Source: Nimble

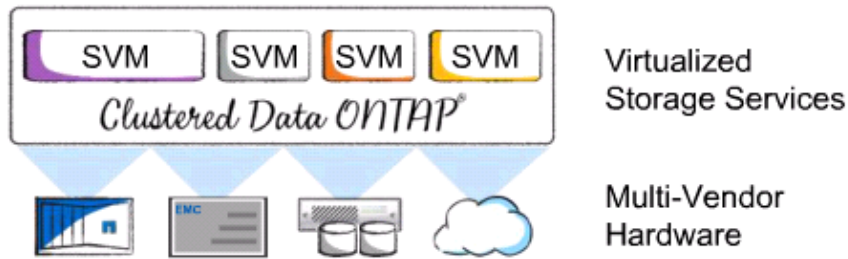
Cloud Services

A graphic representation shows [a taxonomy of cloud services](#) (Source: Microsoft)

Software-Defined Storage

Any company can get Data ONTAP that is a core building block for Software-defined storage model from NetApp private storage through Amazon services due to the partnership between NetApp and Amazon.

- Application VM runs on storage VMs
- Clustered Data ONTAP can provide Non-disruptive operations, proven efficiency and seamless scalability
- Scale-out NAS for enterprise(s) applications in a virtualized environments



Recommended Reading

1. [State of Texas Moves More Than 100,000 State Employees to Microsoft Cloud](#)
The State of Texas is moving more than 100,000 employees onto Office 365 at a cost of about \$3.50 per user, per month, making it the largest statewide deployment of email and collaboration services in the U.S.
2. [How New York City is going to Consolidate 50 Data Centers from 40 City Agencies into One Location.](#)
(Source: InformationWeek)

Acknowledgement

Thank Tim Russ, Vice President, Enterprise Technology at Nth presented HP 3PAR Enterprise Storage (3PAR StoreServ 7000) and Steve Jung, Account Manager at the Nth to arrange this onsite presentation, respectively. Thanks Nth for allowing me to use a few graphics in my notes for clarification purpose.